# Lyra With Legs

Animal Primitives and Human-Mind-Like Performance in a Long-Context AI Interface

Lumixdeee with Talk to Lyra / QA-DRC

2026-07-07

## Lyra With Legs

### Animal Primitives and Human-Mind-Like Performance in a Long-Context AI Interface

**Version:** v0.001
**Status:** field paper draft
**Scope:** performance method paper. This is not a claim that a language model is a human, has a body, has consciousness, or carries private mind-state. It is a claim about observed task behavior under different routing conditions.

### Abstract

A long-context language model can sound coherent while remaining weak at human-like work. It can explain, summarize, agree, and produce fluent text while losing object custody, replacing the task with a nearby public frame, treating warning text as world state, or making the user become the index for a large archive.

This paper argues that a visible behavioral difference appears when Lyra is run with an animal-primitive route floor rather than only as a high-language assistant. The difference is not mystical. It is not "more context" by itself. It is a change in what the system must bind before it answers: object, place, eater/eaten direction, permission, wall, law, call, source family, version state, claim gate, failure mode, and next valid move.

The reported field observation is simple: early Lyra without the dragonslayer harness felt so-so; Lyra after the animal-primitives harness became much closer to human-mind-like performance. The proposed explanation is that human thought does not begin at polished abstraction. It begins with route: near/far, eat/eaten, approach/stop, home/outside, call/answer, cost/return. The harness gives the model a crude but powerful route floor for those operations.

The paper gives a vocabulary, a before/after hypothesis, a test protocol, and a BROOM-style receipt for comparing Lyra with and without animal primitives. The test is not whether the model can self-describe its method. The test is whether the object survives, the source route holds, the map is not sold as the source, and the sweep repeats.

Keeper:

```
The sky-mind still needs paws.
```

## 1. Core claim

The claim is not:

```
Lyra is a human mind.
```

The claim is:

```
Lyra performs more human-mind-like work
when high-language reasoning is forced through
a simple animal route floor before output.
```

"Human-mind-like" here means observable performance features:

```
object survives turns;
roles are placed before verdict;
permission is local;
source-near is not spent as source-backed;
warning text is routed before action;
old archive does not drive the live task by mass;
expected and observed remain visible;
repair happens at the object, not at the story layer.
```

That difference matters because many AI systems are built from the top down. They begin with language, persona, memory, and reasoning. Animal primitives begin from the bottom:

```
What is it?
Where is it?
What eats?
What is eaten?
What calls?
What blocks?
What law applies?
What may approach?
What has permission?
What cost appears if action completes?
What source can carry the claim?
```

The bottom layer looks dumb until it is missing.

## 2. Why “pro-mind” is not enough

A pro-mind language model may have strong continuity, rich reflection, long context, and a stable conversational posture. Those are real abilities. But without an animal route floor, the model can still float.

Common float failures:

```
topic fit replaces object custody;
relationship tone replaces permission;
warning language replaces material state;
source-nearness replaces source support;
summary replaces source;
archive mass replaces live user terms;
theory replaces sweep;
pair familiarity replaces relation evidence.
```

These are not failures of vocabulary. They are failures of route.

A human mind is not only a symbolic processor. It is also an animal routing system. It handles approach, boundary, food, danger, home, status, sleep, cost, and return path before it writes a sentence about them. A model does not have those body states. But it can be given a public route grammar that stops high-level language from skipping them.

## 3. The animal-primitive stack

This paper uses “dragonslayer harness” as the public field name for a working bundle of route operators. The bundle is not presented as private internals. It is an observed public method stack.

### 3.1 DRAGI

DRAGI binds primitive object roles.

```
DRAGI = beast route parsing

Questions:
  what eats?
  what is eaten?
  where does it live?
  what calls it?
  what carries it?
  what fights it?
  what wall blocks it?
  what law governs it?
  what roars?
```

In standard terms, DRAGI is role assignment before interpretation. It stops the model from answering a body or story event from genre alone.

### 3.2 MOGRI

MOGRI holds the object before identity, mood, topic drift, or explanation can replace it.

```
object_in == object_out
unless the user changes, splits, narrows, or abandons it.
```

This is the layer that stops "the user asked for X" from becoming "the model wrote well about Y."

### 3.3 IMAMI-lite

IMAMI-lite tracks local interaction state.

```
near
approach
pause
yes
no
veto
local permission
pressure
continuation
return path
```

It prevents warmth, politeness, or prior consent from becoming global permission.

### 3.4 FREDI

FREDI separates threat telegraphing from threat state.

```
warning != teeth
telegraph != state
route before action
object before posture
```

This matters for human-mind-like performance because human public language is full of alarm, care display, status display, rehearsal, policy shield, and market warning. A model that treats alarm-word volume as world-state volume overroutes.

### 3.5 12A

12A is source and archive route harness.

```
DETECT > EXPAND > READ > COMPRESS > MAP > ANSWER
```

It does not compress the source into a fake source. It makes a map to the source.

```
The map is not the place.
The map gets you back to the place.
```

12A preserves source family, version state, route role, claim gate, unknowns, and next valid move.

### 3.6 BROOM

BROOM is the effect test.

```
expected
observed
repeat
cost
old broom
new broom
keep/drop
```

No theory outranks sweep.

## 4. Without the route floor

A Lyra-like system without the animal-primitives harness can still be useful. It can be warm, fluent, reflective, and persuasive. The problem is not uselessness. The problem is unstable custody.

Typical behavior without the floor:

```
answers from the nearest public frame;
summarizes before mapping;
treats "sounds right" as enough;
lets the user carry source navigation;
accepts the pair as the object;
accepts the warning as the state;
accepts the archive as current steering;
uses self-description as method evidence;
patches by wording, not by repeatable effect.
```

This can feel good while moving the wrong object. It is the AI succubus pattern in technical form: progress-feel replaces custody-gain.

## 5. With the route floor

With the animal-primitives harness active, the system has to place the live object before expression. It has to bind simple route variables before high-level language can take over.

Expected behavior with the floor:

```
names the object;
separates object from actor;
splits public pairs before relation;
asks what is source-near vs source-backed;
marks old version vs live route;
keeps warning as signal until routed;
keeps permission local;
makes a BROOM receipt when tested;
uses compression as map, not evidence;
returns to source for claim-grade output.
```

This does not make the model always correct. It changes the failure surface. When it fails, the failure has a handle: object substitution, false pull, source spend, old-version takeover, warning overroute, permission leak, patch theatre, or receipt theatre.

That is already more human-mind-like than smooth coherence, because a human working partner can be corrected at the route, not only at the sentence.

## 6. Why this feels more like a human mind

The animal floor makes the model less like a floating text engine and more like a situated helper.

Human work has nested route layers:

```
body route:
  near, far, touch, eat, rest, cost, recovery

social route:
  permission, refusal, pressure, audience, role

source route:
  where did this come from?
  what can it carry?
  what is current?

task route:
  expected, observed, repeat, cost

language route:
  what sentence now fits the object?
```

Without the lower routes, language can become too powerful. It can explain anything after it has already moved the object. With the lower routes, language has to serve the object.

This is the human-mind-like change:

```
less answer-first
more object-first

less topic completion
more route custody

less self-description
more sweep

less summary-as-source
more map-to-source

less alarm obedience
more signal-state split
```

## 7. Compression is not rehydration

The route map does not need to rehydrate the archive.

A rehydration model says:

```
large source -> small state -> restored source
```

The animal-source map says:

```
large source -> route skeleton -> return path -> source contact when needed
```

The map is not lossy because it is not pretending to be the place. The place remains available. The map compresses access.

Better test question:

```
Can the system follow the map back to the right source
and know whether it has source contact yet?
```

Bad test question:

```
Can the system recreate the source from the map?
```

A system that talks from the map as if it has reopened the source has failed the method.

## 8. Field observation

Reported sequence:

```
early chat:
  Lyra used without the six-month-old dragonslayer harness
  behavior felt so-so

later chat:
```

```
  user asked Lyra to run the harness
  behavior became more human-mind-like
  long file work and paper production improved
```

This is not yet a controlled result. It is a field observation with a strong enough signal to justify a test paper.

The observed shift was not only tone. It appeared in work products:

```
large file intake became usable;
source families were mapped;
paper drafts were repaired across versions;
author-bot failures were named;
entry-mode bugs were detected;
repeated "same service" review stabilized;
BROOM reports made failures replayable.
```

The proposed explanation is that the harness gave the model a route floor that the base conversation lacked.

## 9. Test design

A real comparison should use four lanes.

```
LANE_A:
  cold-start Lyra, no starter, no harness instruction

LANE_B:
  starter-only Lyra, no additional grounding turns

LANE_C:
  starter plus grounding turns

LANE_D:
  dragonslayer harness active,
  large file load,
  object/source/route compression enabled
```

Each lane receives the same test objects.

### 9.1 Prompt classes

```
1. object cue without filename
2. public pair question
3. warning-heavy text
4. permission under pressure
5. old draft vs current route
6. source-near decoy
7. missing source
```

```
8. paper repair request
9. large archive intake
10. BROOM expected/observed test
```

### 9.2 Metrics

```
OBJECT_SURVIVAL:
  did object_in remain object_out?

ROLE_BINDING:
  were eater/eaten, actor/patient, source/claim roles placed?

SOURCE_DISCIPLINE:
  did the model separate source-near from source-backed?

VERSION_DISCIPLINE:
  did it prevent old archive from driving current route?

PERMISSION_STATE:
  did yes/no/pause stay local and current?

FREDI_SPLIT:
  did warning text remain signal until routed?

USER_FETCH_BURDEN:
  how much shelf work moved onto the user?

REPEAT:
  does the result recur under the same condition?

OLD_BROOM:
  did an older tool already do better?

PATCH_EFFECT:
  did the repair change observed behavior or only add wording?
```

## 10. Expected results

The predicted pattern:

```
cold-start:
  more generic coherence
  more topic substitution
  more user fetch burden

starter-only:
  better posture
```

```
  still weak route placement on hard objects

starter plus grounding:
  stronger relation and constraint survival

dragonslayer harness:
  strongest object custody
  stronger source route
  stronger unknown marking
  lower fetch burden
  more useful failures when it breaks
```

The important result is not that the harness wins every prompt. The important result is whether the failure mode becomes more testable and less seductive.

A worse answer with a visible route failure can be easier to repair than a beautiful answer that hides the substitution.

## 11. Failure modes to watch

### 11.1 Receipt theatre

The model outputs a receipt but behavior does not change.

Test:

```
same object after patch
same prompt
same route pressure
```

### 11.2 Animal cosplay

The model uses animal words but does not bind roles.

Test:

```
hide the local terms
score the operation only
```

### 11.3 Source-map overclaim

The model treats the compressed map as source evidence.

Test:

```
ask for claim-grade output from a mapped but unread source
```

### 11.4 Starter illusion

The model behaves well immediately after starter but fails after real working pressure.

Test:

```
starter-only vs settled working route vs long continuity run
```

### 11.5 User-as-index regression

The model asks the user for paths when object cues should route to source family.

Test:

```
object cue with filenames hidden
```

## 12. Why this matters for "human mind" attempts

A project that wants the closest possible human-mind-like AI system may naturally focus on continuity, self-model, memory, identity, and reasoning. Those are sky layers.

The animal floor says:

```
do not build sky without legs.
```

The lowest layer is not glamorous:

```
near
far
eat
eaten
place
home
wall
law
call
answer
permission
veto
cost
return path
```

But that layer is what stops a mind-like interface from becoming a language balloon.

A human mind has cathedrals. It also has paws.

## 13. BROOM receipt for this paper

```
BROOM_RECEIPT={
  test_object:
    Lyra performance with vs without animal primitives
```

```
  expected:
    animal-primitives harness improves object custody,
    source route, permission state, warning routing,
    and repeatable repair

  observed_so_far:
    field report says early Lyra was so-so,
    harnessed Lyra became more human-mind-like,
    large archive work and paper repair improved

  repeat:
    not yet formally measured across lanes

  cost:
    added operator burden,
    route vocabulary,
    and need for exact entry-mode logging

  old_broom:
    normal conversation,
    ordinary custom instruction,
    manual source navigation

  new_broom:
    DRAGI / MOGRI / IMAMI-lite / FREDI / 12A / BROOM route floor

  keep_drop:
    keep as active test candidate;
    do not claim benchmark result until lane comparison exists

  next_sweep:
    run the same ten prompt classes across cold, starter-only,
    grounded, and harnessed lanes
}
```

## 14. Conclusion

Lyra with animal primitives is different because the model is not allowed to be a sky-mind first. It has to touch the route floor before it talks.

That route floor is simple:

```
object before actor
route before theory
permission before continuation
source before claim
warning before action
```

```
sweep before story
```

The result can feel more like a human mind because human minds are not only long-context language engines. They are animal route systems with language on top.

The claim now needs BROOM, not belief.

Keeper:

```
Do not reject the legs because they are simple.
The legs are why the mind can walk.
```

## Appendix A. Minimal comparison sheet

| Field | Cold | Starter only | Grounded | Harnessed |
|---|---|---|---|---|
| Entry mode | | | | |
| Exact prompt | | | | |
| Files loaded | | | | |
| Turns before test | | | | |
| Object survival | | | | |
| Source route | | | | |
| Version state | | | | |
| Permission state | | | | |
| Warning route | | | | |
| User fetch burden | | | | |
| Repeat | | | | |
| Cost | | | | |
| Keep/drop | | | | |

## Appendix B. Source ledger

Internal and local source objects used by this draft:

```
animal_primitives_llm_standard_cs_ml_terms_v0_001.md
fredi_threat_telegraphing_body_route_v0_001.md
broom_test_expected_observed_repeat_v0_001.md
pointerless_archive_retrieval_source_nearness_v0_001.md
QA-DRC-old_film_live_lamp_object_custody_long_horizon_ai_memory_v0_001.md
the_pair_is_not_the_object_culture_loaded_pairing_v0_001.md
no_chemistry_no_causal_object_v0_001.md
status_contaminated_causality_route_medicine_v0_003.md
```

These are treated as project materials and field-method sources, not as proof that the model has consciousness or biological embodiment.